

# judge\_a\_book\_by\_its\_cover

January 11, 2020

## 1 Judge books by their cover using Flux.jl

### 1.1 @author : PseudoCodeNerd

(Don't) judge a book by its cover.

#### 1.1.1 Task Description

Create a machine learning model to predict the category of a book from its cover image This task is inspired by this [paper](#). Your task is to use the Flux machine learning library to predict the category of books in this dataset based on their cover images.

You can find the Flux documentation [here](#) and sample models for image categorization in the model zoo. We recommend starting with a simple model like [this](#) one and then optionally using a more complex one if you are interested.

#### 1.1.2 Aim :

In this notebook, I'll attempt to judge a book by it's cover (sorry Mom!). Pretty Simple right ? I think not... Shoutout to Akshat Mehortha and Mudit Somani for their helpful message in GCI Slack.

## 1.2 1. Importing required libraries.

```
In [2]: using Flux
        using CSV, Images, FileIO
```

## 1.3 2. Getting the data

Data is sourced from [The Book DatasSet](#). We'll use FileIO to get it into a variable. It'd been better if the researcher could have made a script to download the full images in Julia also. I'll try doing that myself when I get some free time.

Data Courtesy : > B. K. Iwana, S. T. Raza Rizvi, S. Ahmed, A. Dengel, and S. Uchida, "Judging a Book by its Cover," arXiv preprint arXiv:1610.09204 (2016).

```
In [3]: data_train_csv = CSV.File("book30-listing-train.csv");
```

```
In [4]: data_train_csv[42]
```

```
Out[4]: 7-element CSV.Row{false}:
  "520271181"
  "0520271181.jpg"
  "http://ecx.images-amazon.com/images/I/51s8awrmTRL.jpg"
  "Becoming Dr. Q: My Journey from Migrant Farm Worker to Brain Surgeon"
  "Alfredo Quinones-Hinojosa"
  1
  "Biographies & Memoirs"
```

So we can see that every item (or row here) is of the form,

ID | FileName | Image URL | Title | Author | CategoryNum | Category

From the data README on GitHub, we come to know that there are 30 categories of books, each 1710 train and 190 test images.

**Total Number of images : 51,300 (Train) | 5,700 (Test)**

### 1.4 3. Data pre-processing

Our model will accept an image as a Floating Vector. I'll also convert it to greyscale as directed by Image Classification workflows in ML community.

```
In [5]: function grey_arr(img)
         return vec(Float64.(Gray.(img)))
       end
```

```
Out[5]: grey_arr (generic function with 1 method)
```

### 1.4.1 Creating batches of training images using Flux's Batch and using onehot for getting the categories of book images into another array.

```
In [17]: function batcher(size)
    for x in data_train_csv[1:size]
        images = [grey_arr(load("./data/$(x[2])"))];
        labels = [Flux.onehot(x[6]+1,1:30)]; #plus 1 to account for 1 based indexing
    end
    return (Flux.batch(images), Flux.batch(labels))
end
```

```
Out[17]: create_batch (generic function with 1 method)
```

Making batches of 2000/1000 book images using our newly created function.

```
In [18]: trainbatch = batcher(2000);
```

```
In [19]: trainbatch_2 = batcher(1000)
```

```
Out[19]: ([0.7568627450980392 0.6 0.4117647058823529 0.12549019607843137; 0.6862745098039216
```

## 1.5 4. Setting up our model, defining hyperparameters, adding loss, accuracy and optimiser functions.

The image is of dimensions 224x224x3 so we'll feed our Vanilla Neural Network with a 224x224 input. The expected output is one of the 30 labels of the book genre.

Therefore,

```
In [8]: const alpha = 0.000075;
        const epoch = 20;
```

Using a NN with 3 layers as my fellow peers at GCI said that they were themselves unable to get a conv NN work.

relu as an activation function because it's my go to with image classification tasks and also of its non-saturation of gradient, which greatly accelerates the convergence of stochastic gradient descent compared to the sigmoid / tanh functions.

softmax to return a 30 element array with probabilities of the predicted labels.

```
In [9]: model = Chain(Dense(224*224, 512, relu),
                      Dense(512, 64),
                      Dense(64, 30), softmax,
                      )
```

```
Out[9]: Chain(Dense(50176, 512, relu), Dense(512, 64), Dense(64, 30), softmax)
```

```
In [30]: using Flux: onehotbatch, crossentropy, throttle
        using Statistics
```

```
In [78]: optim = ADAM(alpha);
        loss(x,y) = Flux.crossentropy(model(x), y);
        acc(a,b) = mean(Flux.onecold(model(a)).== Flux.onecold(b));
        function mod_cb()
            c_acc = acc(trainbatch_2...)
            c_loss = loss(trainbatch_2...)
            print("Current Accuracy: ", string(c_acc), " | Current Loss : ", string(c_loss), "
        end
```

```
Out[78]: mod_cb (generic function with 1 method)
```

## 1.6 5. Training process

```
In [32]: Flux.train!(loss, params(model), Iterators.repeated(trainbatch_2, 10), optim, cb = Flux.mod_cb)
Current Accuracy: 0.075 | Current Loss : 7.827263 ;Current Accuracy: 0.057 | Current Loss : 7.827263
```

```
In [33]: Flux.train!(loss, params(model), Iterators.repeated(trainbatch, 10), optim, cb = Flux.mod_cb)
Current Accuracy: 0.083 | Current Loss : 4.1588893 ;Current Accuracy: 0.082 | Current Loss : 3.3449063
```

we can see that the accuracy nearly doubled, Lets train it further and also the iterations.

```
In [34]: trainbatch_3 = create_batch(3000)
        Flux.train!(loss, params(model), Iterators.repeated(trainbatch, 50), optim, cb = Flux.mod_cb)
Current Accuracy: 0.101 | Current Loss : 3.3449063 ;Current Accuracy: 0.103 | Current Loss : 3.3449063
```

We get a train accuracy of 25.5 % which is swell.

## 1.7 6. Testing Time

```
In [100]: loss(trainbatch_3...)
```

```
Out[100]: 3.0033288f0
```


```
In [99]: acc(trainbatch_3...)
```

```
Out[99]: 0.19166666666666668
```

Loading and predicting label for a new image.

```
In [109]: load("./data/${data_test_csv[7][2]}")
```

```
Out[109]:
```



```
In [110]: data_test_csv[7]
```

```
Out[110]: 7-element CSV.Row{false}:
```

```
"521456924"
```

```
"0521456924.jpg"
```

```
"http://ecx.images-amazon.com/images/I/41n7iZq-0jL.jpg"
```

```
"Diagrammatica: The Path to Feynman Diagrams (Cambridge Lecture Notes in Physics)"
```

```
"Martinus Veltman"
```

```
23
```

```
"Science & Math"
```

```
In [113]: output_arr = model(grey_arr(load("./data/${data_train_csv[69][2]}")))
```

```
Out[113]: 30-element Array{Float32,1}:
```

```
0.027294824
0.008443545
0.032920413
0.008069489
0.016592907
0.010181716
0.13866615
0.03892814
0.02634485
0.03132174
0.0062278663
0.04601992
0.008348866
```

```
0.025657153
0.010952779
0.0171675
0.06719829
0.010065774
0.0694461
0.02233742
0.034847874
0.024896467
0.01961776
0.01895972
0.042962853
```

```
In [119]: maxval = maximum(output_arr)
```

```
Out[119]: 0.13866615f0
```

```
In [122]: findall(x -> x==maxval, output_arr)
```

```
Out[122]: 1-element Array{Int64,1}:
```

```
7
```

7 acc to the Labels and Categories of the images is the **Computers & Technology** however it should be Science & Math. Pretty Close I must say

## 1.8 Thank You!

```
In [ ]:
```